

RPC:


Sequence of events

1. The client calls the client [stub](#). The call is a local procedure call, with parameters pushed on to the stack in the normal way.
2. The [client stub](#) packs the parameters into a message and makes a system call to send the message. Packing the parameters is called [marshalling](#).
3. The client's local [operating system](#) sends the message from the client machine to the server machine.
4. The local [operating system](#) on the server machine passes the incoming packets to the [server stub](#).
5. The server stub unpacks the parameters from the message. Unpacking the parameters is called [unmarshalling](#).
6. Finally, the server stub calls the server procedure. The reply traces the same steps in the reverse direction.

XML-RPC is a [remote procedure call](#) (RPC) [protocol](#) which uses [XML](#) to encode its calls and [HTTP](#) as a transport mechanism.

SOAP (successor of XML-RPC):

SOAP (abbreviation for **Simple Object Access Protocol**) is a messaging [protocol](#) specification for exchanging structured information in the implementation of [web services](#) in [computer networks](#). Its purpose is to provide [extensibility](#), [neutrality](#) and independence. It uses [XML Information Set](#) for its [message format](#), and relies on [application layer](#) protocols, most often [Hypertext Transfer Protocol](#) (HTTP) or [Simple Mail Transfer Protocol](#) (SMTP), for message negotiation and transmission.



SOAP allows processes running on disparate operating systems (such as [Windows](#) and [Linux](#)) to communicate using [Extensible Markup Language](#) (XML). Since Web protocols like HTTP are installed and running on all operating systems, SOAP allows clients to invoke web services and receive responses independent of language and platforms.

RMI:

In [computing](#), the **Java Remote Method Invocation** (Java




RMI) is a [Java API](#) that performs [remote method invocation](#), the object-oriented equivalent of [remote procedure calls](#) (RPC), with support for direct transfer of [serialized](#) Java classes and [distributed garbage-collection](#).

The original implementation depends on [Java Virtual Machine](#) (JVM) class-representation mechanisms and it thus only supports making calls from one JVM to another. The protocol underlying this Java-only implementation is known as [Java Remote Method Protocol](#) (JRMP). In

order to support code running in a non-JVM context, programmers later developed a [CORBA](#) version.

MOM:

Message-oriented middleware (MOM) is software or hardware infrastructure supporting sending and receiving messages between [distributed systems](#). MOM allows [application modules](#) to be



distributed over heterogeneous platforms and reduces the complexity of developing applications that span multiple [operating systems](#) and [network protocols](#). The [middleware](#) creates a distributed communications layer that insulates the [application developer](#) from the details of the various operating systems and network interfaces. [APIs](#) that extend across diverse platforms and networks are typically provided by MOM.[\[1\]](#)

JMS (a MOM):

The **Java Message Service (JMS) API** is a [Java message-oriented middleware](#) API for sending messages between two or more [clients](#).[\[1\]](#) It is an implementation to handle the [producer–consumer problem](#). JMS is a part of the [Java Platform, Enterprise Edition](#) (Java EE), and was defined by a specification developed at Sun Microsystems, but which has since been guided by the [Java Community Process](#).[\[2\]](#) It is a messaging standard that allows application components based on Java EE to create, send, receive, and read messages. It allows the communication between different components of a [distributed application](#) to be [loosely coupled](#), reliable, and asynchronous.[\[3\]](#)

Messaging is a form of [loosely coupled](#) distributed communication, where in this context the term 'communication' can be understood as an exchange of messages between software components. Message-oriented technologies attempt to relax *tightly coupled* communication (such as [TCP](#) network [sockets](#), [CORBA](#) or [RMI](#)) by the introduction of an intermediary component. This approach allows software components to communicate with each other indirectly. Benefits of this include message senders not needing to have precise knowledge of their receivers.

The advantages of messaging include the ability to integrate heterogeneous platforms, reduce system bottlenecks, increase scalability, and respond more quickly to change.[\[4\]](#)

The JMS API supports two distinct models:

- Point-to-point
- Publish-and-subscribe

COM:

Component Object Model (COM) is a [binary-interface](#) standard for [software components](#) introduced by [Microsoft](#). The essence of COM is a language-neutral way of implementing objects that can be used in environments different from the one in which they were created, even across machine boundaries. For well-authored components, COM allows reuse of objects with no

knowledge of their internal implementation, as it forces component implementers to provide well-defined interfaces that are separated from the implementation. For some applications, COM has been replaced at least to some extent by the Microsoft .NET framework, and support for Web Services through the Windows Communication Foundation (WCF). However, COM objects can be used with all .NET languages through .NET COM Interop. Networked DCOM uses binary proprietary formats, while WCF encourages the use of XML-based SOAP messaging. COM is very similar to other component software interface technologies, such as CORBA and Enterprise JavaBeans.

DCE:

In computing, the **Distributed Computing Environment (DCE)** software system was developed in the early 1990s from the work of the Open Software Foundation (OSF), a consortium (founded in 1988) that included Apollo Computer (part of Hewlett-Packard from 1989), IBM, Digital Equipment Corporation, and others.[1][2] The DCE supplies a framework and a toolkit for developing client/server applications.[3] The framework includes:

- a remote procedure call (RPC) mechanism known as DCE/RPC
- a naming (directory) service
- a time service
- an authentication service
- a distributed file system (DFS) known as DCE/DFS

DCE represented a big step in the direction of standardization of architectures, which had previously been manufacturer-dependent. Like the OSI model, DCE has not seen much success in practical implementation; however, its underlying concepts have had more substantial influence over subsequent efforts.

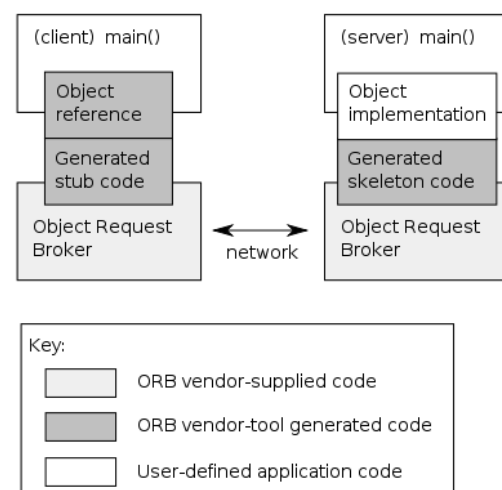
Major components of DCE within every cell are:

1. The **Security Server** that is responsible for authentication
2. The **Cell Directory Server** (CDS) that is the repository of resources and ACLs and
3. The **Distributed Time Server** that provides an accurate clock for proper functioning of the entire cell

One of the major uses of DCE today is Microsoft's DCOM and ODBC systems, which use DCE/RPC (in MSRPC) as their network transport layer.

CORBA:

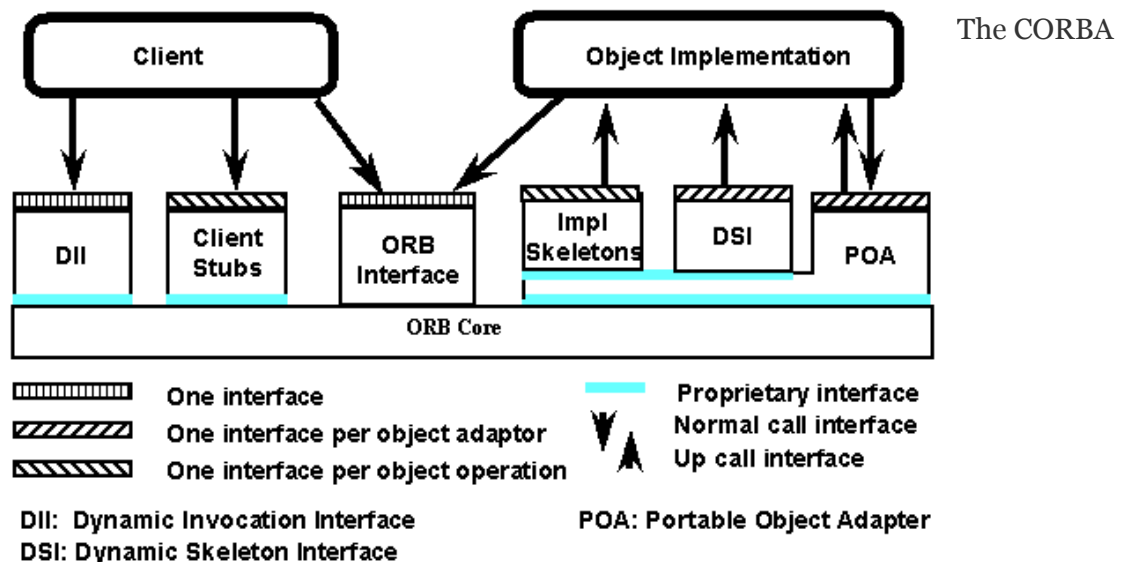
The **Common Object Request Broker Architecture (CORBA)** is a standard defined by the Object Management Group (OMG) designed to facilitate the communication of systems that are deployed on diverse platforms. CORBA enables collaboration between systems on different operating systems, programming languages, and computing hardware. CORBA uses an object-oriented model although the



systems that use the CORBA do not have to be object-oriented. CORBA is an example of the [distributed object](#) paradigm.

CORBA enables communication between software written in different languages and running on different computers. Implementation details from specific operating systems, programming languages, and hardware platforms are all removed from the responsibility of developers who use CORBA.

CORBA uses an [interface definition language](#) (IDL) to specify the interfaces that objects present to the outer world. CORBA then specifies a *mapping* from IDL to a specific implementation language like [C++](#) or [Java](#).



specification dictates there shall be an ORB through which an application would interact with other objects. This is how it is implemented in practice:

1. The application simply initializes the ORB, and accesses an internal *Object Adapter*, which maintains things like reference counting, object (and reference) instantiation policies, and object lifetime policies.
2. The Object Adapter is used to register instances of the *generated code classes*. Generated code classes are the result of compiling the user IDL code, which translates the high-level interface definition into an OS- and language-specific class base for use by the user application. This step is necessary in order to enforce CORBA semantics and provide a clean user process for interfacing with the CORBA infrastructure.

In addition to providing users with a language and a platform-neutral [remote procedure call](#) (RPC) specification, CORBA defines commonly needed services such as transactions and security, events, time, and other domain-specific interface models.

RMI-IIOP is the method that is chosen by Java programmers who want to use the RMI interfaces, but use IIOP as the transport. RMI-IIOP requires that all remote interfaces are defined as Java RMI interfaces. **Java IDL** is an alternative solution, intended for CORBA programmers who want to program in Java to implement objects that are defined in IDL. The general rule that is suggested by Oracle is to use Java IDL when you are using Java to access existing CORBA resources, and RMI-IIOP to export RMI resources to CORBA.

IDL:

An **interface description language** or **interface definition language (IDL)**, is a [specification language](#) used to describe a [software component's application programming interface](#) (API). IDLs describe an interface in a [language-independent](#) way, enabling communication between software components that do not share one language. For example, between those written in [C++](#) and those written in [Java](#).

IDLs are commonly used in [remote procedure call](#) software. In these cases the machines at either end of the *link* may be using different [operating systems](#) and computer languages. IDLs offer a bridge between the two different systems.

Software systems based on IDLs include [Sun's ONC RPC](#), [The Open Group's Distributed Computing Environment](#), [IBM's System Object Model](#), the [Object Management Group's CORBA](#) (which implements OMG IDL, an IDL based on DCE/RPC) and [Data Distribution Service](#), [Mozilla's XPCOM](#), [Microsoft's Microsoft RPC](#) (which evolved into [COM](#) and [DCOM](#)), [Facebook's Thrift](#) and [WSDL](#) for [Web services](#).

Enterprise Java Beans: (wikipedia and <https://stackify.com/enterprise-java-beans/>)

Simply put, an Enterprise Java Bean is a Java class with one or more annotations from the EJB spec which grant the class special powers when running inside of an EJB container.

EJB is one of several [Java APIs](#) for modular construction of [enterprise software](#). EJB is a [server-side software component](#) that [encapsulates business logic](#) of an application. An EJB [web container](#) provides a [runtime environment](#) for web related software components, including [computer security](#), [Java servlet lifecycle management](#), [transaction processing](#), and other [web services](#). The EJB specification is a subset of the [Java EE](#) specification.[1]

The EJB specification provides a standard way to implement the server-side (also called "[back-end](#)") 'business' software typically found in enterprise applications (as opposed to 'front-end' [user interface](#) software). Such software addresses the same types of problem, and solutions to these problems are often repeatedly re-implemented by programmers. Enterprise JavaBeans is intended to handle such common concerns as [persistence](#), [transactional integrity](#) and [security](#) in a standard way, leaving programmers free to concentrate on the particular parts of the enterprise software at hand.

- Session beans – invoked by client

A session bean encapsulates business logic that can be invoked programmatically by a client. The invocation can be done locally by another class in the same JVM or remotely over the network from another JVM. The bean performs the task for the client, abstracting its complexity similar to a web service, for example.

```
@Stateless
public class TestStatelessEjb {
    public String sayHello(String name) {
        return "Hello, " + name + "!";
    }
}
```

- Message driven beans - event-driven

A message-driven bean or MDB is an enterprise bean that allows you to process messages asynchronously. This type of bean normally acts as a JMS message listener, which is similar to an event listener but receives JMS messages instead of events.

They are in many ways similar to a Stateless session bean but they are not invoked by a client. instead, they are event-driven

```
@MessageDriven(mappedName = "jms/TestQueue")
public class TestMessageDrivenBean implements MessageListener {

    @Resource
    MessageDrivenContext messageDrivenContext;

    public void onMessage(Message message) {
        try {
            if (message instanceof TextMessage) {
                TextMessage msg = (TextMessage) message;
                msg.getText();
            }
        } catch (JMSEException e) {
            messageDrivenContext.setRollbackOnly();
        }
    }
}
```

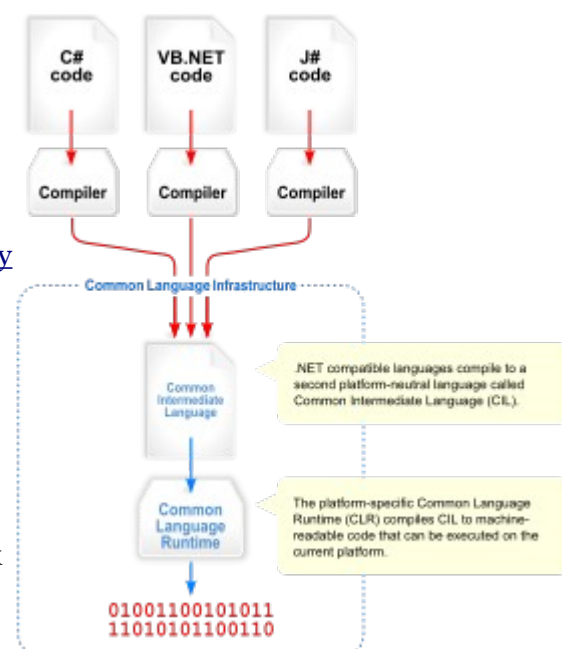
CLI:

The **Common Language Infrastructure (CLI)** is an open [specification](#) (technical standard) developed by [Microsoft](#) and standardized by [ISO\[1\]](#) and [Ecma\[2\]](#) that describes executable code and a runtime environment that allows multiple high-level languages to be used on different computer platforms without being rewritten for specific architectures. This implies it is platform agnostic. The [.NET Framework](#), [.NET Core](#), [Mono](#), [DotGNU](#) and [Portable.NET](#) are implementations of the CLI.

.NET:

.NET Framework (pronounced as "dot net") is a [software framework](#) developed by [Microsoft](#) that runs primarily on [Microsoft Windows](#). It includes a large [class library](#) named as [Framework Class Library](#) (FCL) and provides [language interoperability](#) (each language can use code written in other languages) across several [programming languages](#). Programs written for .NET Framework execute in a [software](#) environment (in contrast to a [hardware](#) environment) named the [Common Language Runtime](#) (CLR). The CLR is an [application virtual machine](#) that provides services such as security, [memory management](#), and [exception handling](#). As such, computer code written using .NET Framework is called "[managed code](#)". FCL and CLR together constitute the .NET Framework.

FCL provides [user interface](#), [data access](#), [database connectivity](#), [cryptography](#), [web application](#) development, numeric [algorithms](#), and [network](#)



[communications](#). Programmers produce software by combining their [source code](#) with .NET Framework and other libraries.

RAID:

RAID (Redundant Array of Inexpensive Disks^[1] or Drives, or Redundant Array of Independent Disks) is a data [storage virtualization](#) technology that combines multiple physical [disk drive](#) components into one or more logical units for the purposes of [data redundancy](#), performance improvement, or both. This was in contrast to the previous concept of highly reliable mainframe disk drives referred to as "single large expensive disk" (SLED).^{[2][3]}

Data is distributed across the drives in one of several ways, referred to as RAID levels, depending on the required level of [redundancy](#) and performance. The different schemes, or data distribution layouts, are named by the word "RAID" followed by a number, for example RAID 0 or RAID 1.

Each scheme, or RAID level, provides a different balance among the key goals: [reliability](#), [availability](#), [performance](#), and [capacity](#).

[RAID 0](#) consists of [striping](#), but no [mirroring](#) or [parity](#).

[RAID 1](#) consists of data mirroring, without parity or striping.


[RAID 2](#) consists of bit-level striping with dedicated [Hamming-code](#) parity.

[RAID 3](#) consists of byte-level striping with dedicated parity.

[RAID 4](#) consists of block-level striping with dedicated parity.

[RAID 5](#) consists of block-level striping with distributed parity.

[RAID 6](#) consists of block-level striping with double distributed parity.



NFS:

is a [distributed file system](#) protocol originally developed by [Sun Microsystems](#) in 1984,^[1] allowing a user on a client [computer](#) to access files over a [computer network](#) much like local storage is accessed. NFS, like many other protocols, builds on the [Open Network Computing Remote Procedure Call](#) (ONC RPC) system. The NFS is an open standard defined in a [Request for Comments](#) (RFC), allowing anyone to implement the protocol.

Assuming a Unix-style scenario in which one machine (the [client](#)) needs access to data stored on another machine (the NFS [server](#)):

1. The server implements NFS [daemon](#) processes, running by default as `nfsd`, to make its data generically available to clients.
2. The server administrator determines what to make available, exporting the names and parameters of [directories](#), typically using the `/etc/exports` configuration file and the `exportfs` command.
3. The server [security](#)-administration ensures that it can recognize and approve validated clients.
4. The server network configuration ensures that appropriate clients can negotiate with it through any [firewall](#) system.
5. The client machine requests access to exported data, typically by issuing a `mount` command. (The client asks the server (`rpcbind`) which port the NFS server is using, the client connects to the NFS server (`nfsd`), `nfsd` passes the request to `mountd`)
6. If all goes well, users on the client machine can then view and interact with mounted [filesystems](#) on the server within the parameters permitted.

Note that automation of the NFS mounting process may take place — perhaps using `/etc/fstab` and/or [automounting](#) facilities.

VFS:

A **Virtual File System (VFS)** or **virtual filesystem switch** is an abstract layer on top of a more concrete [file system](#). The purpose of a VFS is to allow client applications to access different types of concrete file systems in a uniform way. A VFS can, for example, be used to access [local](#) and network storage devices transparently without the client application noticing the difference. It can be used to bridge the differences in [Windows](#), [classic Mac OS](#)/[macOS](#) and [Unix](#) filesystems, so that applications can access files on local file systems of those types without having to know what type of file system they are accessing.

ACID:

In [computer science](#), **ACID** ([Atomicity](#), [Consistency](#), [Isolation](#), [Durability](#)) is a set of properties of [database transactions](#) intended to guarantee validity even in the event of errors, power failures, etc. In the context of [databases](#), a sequence of database operations that satisfies the ACID properties (and these can be perceived as a single logical operation on the data) is called a transaction.

TP-Monitor:

A **teleprocessing monitor** (also, **Transaction Processing Monitor** or **TP Monitor**) is a control [program](#) that monitors the transfer of data between multiple local and remote [terminals](#) to ensure that the transaction processes completely or, if an error occurs, to take appropriate actions. [\[1\]](#)

Distributed Transaction Processing:

A **distributed transaction** is a [database transaction](#) in which two or more network hosts are involved. Usually, hosts provide **transactional resources**, while the **transaction manager** is responsible for creating and managing a global transaction that encompasses all operations against such resources. Distributed transactions, as any other [transactions](#), must have all four [ACID](#) ([atomicity](#), [consistency](#), [isolation](#), [durability](#)) properties, where atomicity guarantees all-or-nothing outcomes for the unit of work (operations bundle).

JBDC, SQLJ, JDO, ODMG, EJB QL

(https://www.service-architecture.com/articles/database/comparison_of_dbms_standards.html)

Feature	SQL-92	JDBC	SQLJ	SQL:1999	ODMG 3.0	JDO
Model	relational model	relational model	Parts 0 & 1: relational model Part 2: SQL:1999 object model (more)	SQL:1999 object model (more)	Java, C++, and Smalltalk object models enhanced for transparent persistence The model with the transparent persistence enhancements is a superset of the OMG Common Object Model	Java object model enhanced for transparent persistence
Data Definition Language	SQL	SQL	SQL	SQL	Object Definition Language (ODL) which is a superset of the OMG Interface Definition Language (IDL)	Java & XML
Query Language	Embedded SQL, Dynamic SQL, and Call-level interface	Call-level interface for SQL	Embedded SQL	Embedded SQL, Dynamic SQL, and Call-level interface	Object Query Language (OQL) which is based on SQL-92	JDO Query Language (JDOQL) (more)
Data Manipulation Language	Embedded SQL, Dynamic SQL, and Call-level interface	Call-level interface for SQL and Java	Embedded SQL and Java	Embedded SQL, Dynamic SQL, and Call-level interface	Java, C++, or Smalltalk	Java
Specification Based On	SQL-89	Java & SQL-92	Java, JDBC & SQL:1999	SQL-92	OMG Common Object Model, OMG IDL, SQL-92, Java, C++, and Smalltalk	Java & XML